

VALIDATION PLAN

OPERATIONAL QUALIFICATION

COMMUNITY SCRIPT SUITE

JR Validated Environment — v2.6.0 target

Document ID	JR-VP-002
Version	4.0
Date	2026-10-01
Status	Approved
Author	Joep Rous
System	JR Validated Environment v5.0.0 (target)

Version	Date	Author	Description
1.0	2026-03-15	Joep Rous	Initial OQ validation plan — 116 test cases, 22 R + 2 Python community scripts.
2.0	2026-03-29	Joep Rous	Phase 9: added 25 numeric correctness TCs for high-risk sample-size and tolerance-interval scripts; added DOE scripts (UR-023) and core infrastructure tests (UR-024); added verify_attr_known.csv; total 177 test cases.
3.0	2026-04-19	Joep Rous	Added jrc_verify_discrete (UR-025): exact Clopper-Pearson pass/fail verification with margin reporting. 8 new test cases (TC-VER-DISC-001..008) including independent numerical verification via pure-Python bisection. Total 185 test cases.
4.0	2026-10-01	Joep Rous	jrscripsts v5.0.0, code review 2026-10: exact binomial sample sizes (TC-DISC-001/006..009, TC-FAT-006/007 updated), signed sample k-factor (TC-ATTR-009, TC-ATTRCK-006, TC-ATTRCI-005), Box-Cox acceptance (TC-ATTR-010),

			TOST z(1-beta/2) (TC-EQUIV-005/006), two-sided Grubbs (TC-OUT-005), normality message (TC-NORM-006), Bland-Altman row check (TC-BA-006), fatigue beta table (TC-FAT-008), 6-sigma GRR design (TC-GRR-005), DoE factor levels, aliasing and saturated models (TC-DOE-DES-019, TC-DOE-ANA-016..018). Total 198 test cases.
--	--	--	--

1. Purpose

This Validation Plan defines the Operational Qualification (OQ) test strategy and test cases for the community script suite of the JR Validated Environment. It provides objective evidence that each community script produces correct, reproducible outputs for its intended use-case and fails predictably on invalid inputs, in accordance with 21 CFR Part 820.70(i) and ISO 13485 §7.5.6.

Version 2.0 of this plan extends v1.0 by adding numeric correctness test cases (Phase 9) for the nine sample-size and tolerance-interval scripts identified as highest risk under 21 CFR 820.70(i): a wrong sample size or a false / tolerance-interval verdict directly invalidates a design-verification study conclusion. Each new test case asserts a computed numeric value against an independently derived expected result (closed-form formula or published K-factor table) — not by running the script.

Version 3.0 adds jrc_verify_discrete (UR-025) — a Clopper-Pearson exact binomial CI tool for discrete pass/fail verification assessment — with 8 new test cases (TC-VER-DISC-001..008), bringing the total to 185 test cases across 23 R scripts and 2 Python scripts.

2. Scope

2.1 In Scope

- All 23 community R scripts (jrc_ss_discrete, jrc_ss_discrete_ci, jrc_ss_attr, jrc_ss_attr_check, jrc_ss_attr_ci, jrc_ss_sigma, jrc_ss_fatigue, jrc_ss_paired, jrc_ss_equivalence, jrc_msa_grr_design, jrc_normality, jrc_outliers, jrc_capability, jrc_descriptive, jrc_bland_altman, jrc_weibull, jrc_verify_attr, jrc_verify_discrete, jrc_gen_normal, jrc_gen_lognormal, jrc_gen_uniform, jrc_doe_generate, jrc_doe_analyse)
- 2 community Python scripts (jrc_convert_csv.py, jrc_convert_txt.py)
- Core infrastructure scripts (jrc_R_hello.R, jrc_py_hello.py, integrity enforcement, bypass protection)
- Correct output computation including numeric value assertions for high-risk scripts
- Error handling and input validation for all scripts
- File I/O behaviour and output file creation
- RENV_PATHS_ROOT enforcement (R scripts exit if run outside jrrun)

2.2 Out of Scope

- Infrastructure scripts (bin/jrrun, admin/* scripts)
- Post-Market Surveillance (PQ) or user acceptance testing (UAT)
- Language interpreter validation (R, Python)
- Third-party package validation (renv, tidyverse, etc.)
- Non-normal distribution code paths for jrc_ss_attr and jrc_verify_attr (Box-Cox transform path; deferred to Phase 10)

3. Regulatory Basis

Reference	Applicability
21 CFR Part 820.70(i)	FDA Quality System Regulation — software validation for production and process controls
21 CFR Part 11	Electronic records and signatures
FDA GPSV (2022)	General Principles of Software Validation — risk-based testing, repeatability
ISO 13485:2016 §7.5.6	Validation of processes for production and service provision
ASTM E2537	Guide for application of continuous quality verification to pharmaceutical and biopharmaceutical manufacturing

4. Definitions

Term	Definition
IQ	Installation Qualification — verifies that the software is installed correctly
OQ	Operational Qualification — verifies that the software operates as intended across its operating range
PQ	Performance Qualification — verifies the software performs correctly under real-world conditions
UR	User Requirement — a testable statement of what the system must do
TC	Test Case — a specific test with defined inputs, expected outputs, and pass/fail criteria
jrrun	The validated execution wrapper (bin/jrrun) that enforces integrity, environment, and logging
renv	R package manager used to pin and isolate package versions
venv	Python virtual environment used to pin and isolate package versions
CSV	Comma-separated values file — standard input format for all data scripts
TOST	Two One-Sided Tests — the statistical method used for equivalence testing
MLE	Maximum Likelihood Estimation — used for Weibull parameter fitting
LoA	Limits of Agreement — output of Bland-Altman analysis
TI	Tolerance Interval — a statistical interval containing at least proportion P of a population with confidence C
LSL/USL	Lower/Upper Specification Limit — the engineering limits against which the tolerance interval is compared
K-factor	Tolerance interval multiplier; K1 = one-sided, K2 = two-sided

5. System Description

The JR Validated Environment is a framework for running validated R and Python analysis scripts in a controlled, reproducible environment for FDA medical device development. Package versions are pinned and installed from a local repository. Every run is integrity-checked and logged.

Component	Details
Operating system	macOS (arm64) / macOS 15.7.3 Darwin 24.6.0
R version	4.5.0
Python version	3.11.9
Entry point	bin/jrrun
R environment manager	renv (version pinned in admin/renv.lock)
Python environment mgr	venv (packages in admin/python_requirements.txt)
Integrity mechanism	SHA-256 checksum of all tracked files (admin/project_integrity.sha256)
Logging	~/jrscript/<PROJECT_ID>/run.log (each run)
OQ test runner	pytest 8.3.4 (frozen, separate venv)
OQ test directory	oq/
OQ data directory	oq/data/

6. Risk Assessment

Risk drives test case priority. The following risk factors are considered:

Risk Factor	Affected Scripts	Mitigation
Incorrect sample size calculation	jrc_ss_discrete, jrc_ss_attr, jrc_ss_sigma, jrc_ss_fatigue, jrc_ss_paired, jrc_ss_equivalence	Numeric assertions against independently derived expected values (closed-form chi-squared, K-factor table, normal quantiles); cross-script consistency checks
False   verdict in tolerance interval check	jrc_verify_attr, jrc_ss_attr_check	Boundary test cases using analytically constructed dataset (verify_attr_known.csv) with exact mean and SD; LSL chosen to straddle $K1(30,0.95,0.95)=2.220$
Silent computational error (correct exit code, wrong value)	All statistical scripts	Numeric output assertions with tolerance ± 0.050 for all high-risk scripts; independent derivation basis documented per test case
Incorrect Clopper-Pearson proportion	jrc_ss_discrete_ci, jrc_ss_attr_ci	Expected proportion derived from beta distribution CDF; assertions use ± 0.0001 tolerance
False PASS/FAIL verdict in discrete verification	jrc_verify_discrete	TC-VER-DISC-004 independently verifies the upper CI bound via pure-Python bisection on the Binomial CDF; TC-VER-DISC-005 uses the exact formula $1-(1-C)^{(1/N)}$ for $f=0$ — both computed without R. Boundary cases tested by TC-VER-DISC-001 (PASS) and TC-VER-DISC-002 (FAIL).
File I/O failure	All scripts	Explicit test cases for nonexistent files, missing arguments, and mismatched row counts
Environment bypass (running R script outside jrrun)	All R scripts	RENV_PATHS_ROOT enforcement tested in TC-CORE-OQ-004
Project integrity tampering	Core infrastructure	SHA-256 integrity check tested in TC-CORE-OQ-003

7. User Requirements

User requirements are grouped by functional area. Each requirement maps to one or more test cases in Section 9. The traceability matrix in Section 10 provides the complete mapping.

UR	Requirement Statement	Test Cases
UR-001	jrc_ss_discrete shall compute the required sample size for a zero-failure (or f-failure) demonstration test using the exact binomial minimum (smallest n with $\text{pbinom}(f, n, 1-P) \leq 1-C$); the result shall equal the independently derived value.	TC-DISC-001..009
UR-002	jrc_ss_discrete_ci shall compute the exact Clopper-Pearson confidence interval on the demonstrated proportion given n and f; the result shall match the independently derived value to within ± 0.0001 .	TC-DISCICI-001..006
UR-003	jrc_ss_attr shall compute the required sample size for an attribute tolerance interval study and shall check whether a planned sample size is sufficient using the K-factor method.	TC-ATTR-001..010
UR-004	jrc_ss_attr_check shall evaluate whether a planned sample size n passes or fails against $K1(n,P,C)$ for the given dataset; the pass/fail boundary shall be correctly determined around $K1(30,0.95,0.95)=2.220$.	TC-ATTRCK-001..006
UR-005	jrc_ss_attr_ci shall compute the achieved proportion for a given dataset and sample size using the K-factor method at the specified confidence level.	TC-ATTRCI-001..005
UR-006	jrc_ss_sigma shall compute the required sample size for a sigma (continuous) pilot study using the normal approximation formula for both 1-sided and 2-sided specifications.	TC-SIGMA-001..007
UR-007	jrc_ss_fatigue shall compute the required sample size for a fatigue demonstration test accounting for Weibull-adjusted acceleration factors.	TC-FAT-001..008
UR-008	jrc_ss_paired shall compute the required sample size for a paired t-test study for both 1-sided and 2-sided specifications.	TC-PAIRED-001..007
UR-009	jrc_ss_equivalence shall compute the required sample size for a TOST equivalence study; results shall equal the 1-sided paired formula for the same inputs.	TC-EQUIV-001..006
UR-010	jrc_msa_grr_design shall compute the required sample size for a Gauge R&R study for both process-based and tolerance-based approaches.	TC-GRR-001..005
UR-011	jrc_normality shall perform a normality test and report the test statistic, p-value, and a pass/fail conclusion at $\alpha=0.05$.	TC-NORM-001..006
UR-012	jrc_outliers shall detect outliers using the Grubbs test and report identified outlier positions.	TC-OUT-001..005
UR-013	jrc_capability shall compute Cp, Cpk, Pp, Ppk and produce a capability chart.	TC-CAP-001..004
UR-014	jrc_descriptive shall compute mean, SD, median, IQR, skewness, kurtosis, and min/max for a dataset.	TC-DESC-001..004
UR-015	jrc_bland_altman shall compute the bias, limits of agreement, and proportional bias test for two measurement methods.	TC-BA-001..006
UR-016	jrc_weibull shall fit a Weibull distribution to time-to-failure data, report shape (β), scale (η), and B-life values, and handle censored data.	TC-WEIB-001..006
UR-017	jrc_verify_attr shall compute a statistical tolerance interval for a dataset and compare it against the specified limits, producing a binding  /  verdict; the K-factor shall match the independently derived value $K1(30,0.95,0.95)=2.220$.	TC-VER-001..011
UR-018	jrc_gen_normal, jrc_gen_lognormal, jrc_gen_sqrt, jrc_gen_boxcox, and jrc_gen_uniform shall generate synthetic datasets with the specified parameters, reproducibly with a fixed seed.	TC-GEN-*
UR-019	jrc_convert_csv.py shall convert a multi-column CSV file to the standard id/value two-column format for a specified column.	TC-CCSV-001..008
UR-020	jrc_convert_txt.py shall convert a plain-text (single or multi-column) file to the standard id/value two-column CSV format.	TC-CTXT-001..007
UR-021	All scripts shall validate their input arguments and exit with a non-zero code if required arguments are missing or invalid, printing a Usage message.	All error-path TCs
UR-022	All R scripts shall exit with a non-zero code if RENV_PATHS_ROOT is not set (i.e., if invoked outside jrrun).	TC-CORE-OQ-004
UR-023	jrc_doe_design shall produce a design of experiments matrix (full factorial, fractional factorial, Plackett-Burman) and jrc_doe_analyse shall analyse the results and identify significant factors.	TC-DOE-DES-001..019, TC-DOE-ANA-001..018
UR-024	Core infrastructure scripts (jrc_R_hello.R, jrc_py_hello.py) shall execute	TC-CORE-IQ-001, TC-CORE-OQ-

	successfully; project integrity verification shall detect file tampering; Python bypass protection shall function equivalently to R bypass protection.	001..005
UR-025	jrc_verify_discrete shall compute the upper one-sided Clopper-Pearson exact binomial confidence bound on the failure rate given N, f, P, and C; compare the bound to the allowable failure rate (1-P); report margin in percentage points; and produce a binding  /  PASS/FAIL verdict. The computed bound shall match the independently derived Clopper-Pearson value to within ± 0.0005 (TC-VER-DISC-004, TC-VER-DISC-005).	TC-VER-DISC-001..008

8. Test Environment

All tests run inside the jrrun validated environment. The OQ test runner uses a separate pytest venv that does not share packages with the analysis venv.

Item	Specification
Host	Joeeps-MacBook-Air.local (Apple M-series, arm64)
OS	macOS 15.x (Darwin 24.6.0, aarch64)
R version	4.5.0 (as installed by admin_install_R)
Python version	3.11.9 (system Python)
pytest version	8.3.4 (frozen in oq/requirements.txt)
OQ venv	~/venvs/MyProject_oq/ (or system python3 with pytest installed)
OQ requirements	pytest==8.3.4 (oq/requirements.txt)
Test runner command	python3 -m pytest oq/
Test data directory	oq/data/ (all synthetic datasets committed to repository)
Working directory	oq/data/ (set by conftest.py for all script invocations)
Script invocation	Via bin/jrrun; all scripts integrity-checked before execution

9. Test Data

All test datasets are synthetic, generated with fixed seeds for reproducibility. No proprietary or clinical data is used. All files are committed to oq/data/.

File	Description	Generation method
normal_n30_mean10_sd1_seed42.csv	30 values, N(10,1), seed 42	R: rnorm(30,10,1) seed 42
skewed_n30_lognormal_seed42.csv	30 log-normal values (skewness ≈ 0.52), seed 42	R: exp(rnorm(30,0,0.5)) seed 42
outlier_n30_seed42.csv	Normal(10,1), row 15 replaced by 5-sigma outlier	R: rnorm(30,10,1) seed 42; row 15 = 15.0
bland_altman_method1_seed42.csv	25 values, N(10,2), seed 42	R: rnorm(25,10,2) seed 42
bland_altman_method2_seed42.csv	Method1 + N(0,0.2) proportional bias, seed 99	R: method1 * (1 + rnorm(25,0,0.02)) seed 99
method1_short.csv	First 10 rows of method1 (mismatch test)	R: head(method1, 10)
weibull_n20_seed42.csv	Weibull(shape=2, scale=1000), 15 failures / 5 censored	R: rweibull(20,2,1000) seed 42
all_censored.csv	20 rows, status=0 (all censored — error path)	Manual construction
neg_times.csv	20 rows, row 1 cycles = -50 (negative time — error path)	Manual construction
bad_status.csv	20 rows, status includes value 2 (invalid — error path)	Manual construction
convert_multicolumn.txt	Tab-delimited, 3 headers (id/ForceN/Temp), 20 rows	Manual construction
convert_singlecolumn.txt	200 numeric values, one per line	R: rnorm(200,10,2) seed 7

verify_attr_known.csv	30 rows: 28 × value=10.0000, row 29 = 13.8079, row 30 = 6.1921. Exact: mean=10.0000, sd=1.0000, skewness=0 (symmetric). Designed so $K1(30,0.95,0.95)=2.220$ gives $TL=7.780$ — used for tolerance-interval numeric assertions.	Algebraic construction: $28 \times 10.0 + \text{symmetric outliers} \pm \sqrt{14.5} \approx \pm 3.8079$
-----------------------	---	---

10. Phase 9 Numeric Test Cases (new in v2.0)

The following test cases were added in v2.0 to assert numerically correct values for the nine highest-risk scripts. Each expected value is derived independently from first principles or published K-factor tables (Hahn & Meeker 2017), not by running the script. Existing functional test cases (TC-DISC-001..005 etc.) are unchanged from v1.0.

10.1 Sample Size — Discrete (jrc_ss_discrete)

Exact binomial minimum (since v1.1, code review 2026-10, COR-18): the smallest n with $\text{pbinom}(f, n, 1-P) \leq 1-C$ (equivalently $1 - \text{qbeta}(C, f+1, n-f) \geq P$). For $f=0$: $n = \text{ceiling}(\ln(1-C) / \ln(P))$. The chi-squared formula used up to v1.0 overestimated n by 0–3.

TC-DISC-006

Description: $P=0.99, C=0.95, f=0 \rightarrow n=299$

Tests: UR-001

Command: `jrrun jrc_ss_discrete.R 0.99 0.95`

Rationale: $\text{ceiling}(\ln 0.05 / \ln 0.99) = \text{ceiling}(298.07) = 299$; $0.99^{299} = 0.04954 \leq 0.05 < 0.99^{298} = 0.05004$

Pass criterion: n at $f=0$ equals 299

TC-DISC-007

Description: $P=0.99, C=0.95, f=1 \rightarrow n=473$

Tests: UR-001

Command: `jrrun jrc_ss_discrete.R 0.99 0.95`

Rationale: Pure Python binomial CDF: $P(X \leq 1 \mid 472, 0.01) = 0.05021 > 0.05$; $P(X \leq 1 \mid 473, 0.01) = 0.04980 \rightarrow 473$

Pass criterion: n at $f=1$ equals 473

TC-DISC-008

Description: $P=0.95, C=0.90, f=0 \rightarrow n=45$

Tests: UR-001

Command: `jrrun jrc_ss_discrete.R 0.95 0.90`

Rationale: $\text{ceiling}(\ln 0.10 / \ln 0.95) = \text{ceiling}(44.89) = 45$; $0.95^{45} = 0.09944 \leq 0.10 < 0.95^{44} = 0.10467$

Pass criterion: n at $f=0$ equals 45

TC-DISC-009

Description: $P=0.99, C=0.99, f=0 \rightarrow n=459$

Tests: UR-001

Command: `jrrun jrc_ss_discrete.R 0.99 0.99`

Rationale: $\text{ceiling}(\ln 0.01 / \ln 0.99) = \text{ceiling}(458.21) = 459$; $0.99^{459} = 0.00992 \leq 0.01 < 0.99^{458} = 0.01002$

Pass criterion: n at $f=0$ equals 459

10.2 Sample Size — Discrete CI (jrc_ss_discrete_ci)

Formula: $\text{proportion} = 1 - \text{qbeta}(1-C, f+1, n-f)$. For $f=0$: $\text{proportion} = 1 - (1-C)^{(1/n)}$.

TC-DISCICI-005

Description: C=0.95, n=300, f=0 → proportion ≈ 0.9900

Tests: UR-002

Command: `jrrun jrc_ss_discrete_ci.R 0.95 300`

Rationale: $p = 1 - (1 - 0.95)^{(1/300)} = 1 - 0.05^{(1/300)} = 1 - 0.990011... \approx 0.9900$

Pass criterion: Extracted proportion ≈ 0.9900 ± 0.0001

TC-DISCICI-006

Description: C=0.95, n=22, f=0 → proportion ≈ 0.8729

Tests: UR-002

Command: `jrrun jrc_ss_discrete_ci.R 0.95 22`

Rationale: $p = 1 - 0.05^{(1/22)} = 1 - 0.8729... \approx 0.8729$ (consistent with TC-DISC-008 inverse)

Pass criterion: Extracted proportion ≈ 0.8729 ± 0.0001

10.3 Sample Size — Attribute (jrc_ss_attr, jrc_ss_attr_check, jrc_ss_attr_ci)

Dataset `verify_attr_known.csv`: mean=10.0000, sd=1.0000, skewness=0. $K1(30, P=0.95, C=0.95) = 2.220$ (NCT distribution, computed by `nct.ppf(0.95, df=29, nc=sqrt(30)*qnorm(0.95))/sqrt(30)`).

TC-ATTR-008

Description: `verify_attr_known.csv`, LSL=7.0, P=0.95, C=0.95 → required n in range [3, 10]

Tests: UR-003

Command: `jrrun jrc_ss_attr.R 0.95 0.95 oq/data/verify_attr_known.csv value 7.0 -`

Rationale: $k_{\text{sample}} = (10 - 7)/1 = 3.000$; $K1(3) \approx 7.655 > 3$, $K1(10) \approx 2.355 < 3$; so $n \in [3, 10]$

Pass criterion: Extracted required n satisfies $3 \leq n \leq 10$, and 'sufficient' appears in output

TC-ATTRCK-004

Description: `verify_attr_known.csv`, LSL=7.0, planned_N=30 → PASS

Tests: UR-004

Command: `jrrun jrc_ss_attr_check.R 0.95 0.95 oq/data/verify_attr_known.csv value 7.0 - 30`

Rationale: $k_{\text{sample}} = 3.000 > K1(30, 0.95, 0.95) = 2.220$ → sample k-factor exceeds threshold → PASS

Pass criterion: 'pass' in output

TC-ATTRCK-005

Description: `verify_attr_known.csv`, LSL=8.5, planned_N=30 → FAIL

Tests: UR-004

Command: `jrrun jrc_ss_attr_check.R 0.95 0.95 oq/data/verify_attr_known.csv value 8.5 - 30`

Rationale: $k_{\text{sample}} = (10 - 8.5)/1 = 1.500 < K1(30, 0.95, 0.95) = 2.220$ → sample k-factor below threshold → FAIL

Pass criterion: 'fail' in output

TC-ATTRCI-004

Description: `verify_attr_known.csv`, C=0.95, LSL=7.0 → achieved proportion ≈ 0.989

Tests: UR-005

Command: `jrrun jrc_ss_attr_ci.R 0.95 oq/data/verify_attr_known.csv value 7.0 -`

Rationale: $k_{\text{sample}} = 3.000$; bisection finds P such that $K1(30, P, 0.95) = 3.000$; since $K1(30, 0.95, 0.95) = 2.220 < 3.000$, achieved P ≈ 0.989

Pass criterion: Extracted proportion in (0.985, 1.000)

Mean on the wrong side of a 1-sided spec (code review 2026-10, COR-11): with LSL=12.0 the signed distance is $k_{\text{sample}} = (10 - 12)/1 = -2$. No tolerance bound can clear the limit, so no sample size, PASS or achieved proportion may be reported.

TC-ATTR-009

Description: verify_attr_known.csv, LSL=12.0 (mean below LSL) → no sample size reported

Tests: UR-003

Command: jrrun jrc_ss_attr.R 0.95 0.95 oq/data/verify_attr_known.csv value 12.0 -

Rationale: $k_{\text{sample}} = (10-12)/1 = -2 \leq 0 \rightarrow$ requirement cannot be demonstrated for any N

Pass criterion: Exit 0; 'no sample size can demonstrate this requirement' in output; no required n; 'sufficient' absent

TC-ATTRCK-006

Description: verify_attr_known.csv, LSL=12.0, planned_N=30 → FAIL

Tests: UR-004

Command: jrrun jrc_ss_attr_check.R 0.95 0.95 oq/data/verify_attr_known.csv value 12.0 - 30

Rationale: $k_{\text{sample}} = -2 < K1(30,0.95,0.95)$ for any N → FAIL

Pass criterion: Exit 0; 'FAIL' in output and 'PASS' absent; k-factor from pilot sample = -2

TC-ATTRCI-005

Description: verify_attr_known.csv, C=0.95, LSL=12.0 → no achieved proportion reported

Tests: UR-005

Command: jrrun jrc_ss_attr_ci.R 0.95 0.95 oq/data/verify_attr_known.csv value 12.0 -

Rationale: $k_{\text{sample}} = -2 \leq 0 \rightarrow$ no proportion P satisfies $K1(30,P,0.95) \leq k_{\text{sample}}$

Pass criterion: Exit 0; 'No conforming proportion can be demonstrated' in output; no proportion value

TC-ATTR-010

Description: attr_bimodal_boxcox_fails.csv, LSL=8.0 → Box-Cox rejected, no sample size

Tests: UR-003

Command: jrrun jrc_ss_attr.R 0.95 0.95 oq/data/attr_bimodal_boxcox_fails.csv value 8.0 -

Rationale: Skewed bimodal mixture: Box-Cox lowers |skew| 1.028 → 0.653 but Shapiro-Wilk on the transformed data gives $p = 5e-5 \leq 0.01$, so the transformation is not accepted (code review 2026-10, COR-15)

Pass criterion: Exit 0; 'Could not compute sample size' in output; 'Box-Cox transformation accepted' absent

10.4 Sample Size — Sigma (jrc_ss_sigma)

Formula: $n = \text{ceiling}(((z_{\alpha} + z_{\beta})/\text{precision})^2) + 1$. 1-sided: $z_{\alpha} = \text{qnorm}(C)$. 2-sided: $z_{\alpha} = \text{qnorm}((1+C)/2)$.

TC-SIGMA-005

Description: precision=1.5, 1-sided spec (spec1=9.0), power=0.90, C=0.95 → n=5

Tests: UR-006

Command: jrrun jrc_ss_sigma.R 1.5 oq/data/normal_n30_mean10_sd1_seed42.csv value 9.0 -

Rationale: $z_{\alpha} = \text{qnorm}(0.95) = 1.6449$; $z_{\beta} = \text{qnorm}(0.90) = 1.2816$; $n = \text{ceiling}((2.926/1.5)^2) + 1 = \text{ceiling}(3.80) + 1 = 5$

Pass criterion: n at power=0.90, C=0.95 equals 5

TC-SIGMA-006

Description: precision=1.0, 2-sided spec, power=0.95, C=0.95 → n=14; output mentions 'N >= 14'

Tests: UR-006

Command: jrrun jrc_ss_sigma.R 1.0 oq/data/normal_n30_mean10_sd1_seed42.csv value 9.0 11.0

Rationale: $z_{\alpha} = \text{qnorm}(0.975) = 1.9600$; $z_{\beta} = \text{qnorm}(0.95) = 1.6449$; $n = \text{ceiling}((3.605/1.0)^2) + 1 = \text{ceiling}(13.0) + 1 = 14$

Pass criterion: n at power=0.95, C=0.95 equals 14 AND 'N >= 14' in output

TC-SIGMA-007

Description: precision=2.0, 1-sided spec, power=0.90, C=0.90 → n=3

Tests: UR-006

Command: `jrrun jrc_ss_sigma.R 2.0 oq/data/normal_n30_mean10_sd1_seed42.csv value 9.0 -`

Rationale: $z_{\alpha} = \text{qnorm}(0.90) = 1.2816$; $z_{\beta} = \text{qnorm}(0.90) = 1.2816$; $n = \text{ceiling}((2.563/2.0)^2) + 1 = \text{ceiling}(1.640) + 1 = 3$

Pass criterion: n at power=0.90, C=0.90 equals 3

10.5 Sample Size — Fatigue (jrc_ss_fatigue)

Weibull-adjusted formula: $p_{\text{eff}} = 1 - R^{AF\beta}$, then the same exact binomial minimum as jrc_ss_discrete (since v1.1, COR-19).

TC-FAT-006

Description: R=0.90, C=0.95, $\beta=2$, AF=1.0 $\rightarrow$ n=29 (equals jrc_ss_discrete P=0.90, C=0.95)

Tests: UR-007

Command: `jrrun jrc_ss_fatigue.R 0.90 0.95 2 1.0`

Rationale: AF=1.0: $p_{\text{eff}} = 0.10$; $\text{ceiling}(\ln 0.05 / \ln 0.90) = \text{ceiling}(28.43) = 29$ ($0.9^{29} = 0.0471$, $0.9^{28} = 0.0523$)

Pass criterion: n at f=0 equals 29

TC-FAT-007

Description: R=0.90, C=0.95, $\beta=2$, AF=2.0 $\rightarrow$ n=8

Tests: UR-007

Command: `jrrun jrc_ss_fatigue.R 0.90 0.95 2 2.0`

Rationale: AF=2, $\beta=2$: $R_{\text{test}} = 0.90^4 = 0.6561$; $n = \text{ceiling}(\ln 0.05 / \ln 0.6561) = \text{ceiling}(7.11) = 8$ ($0.6561^8 = 0.0343$, $0.6561^7 = 0.0523$)

Pass criterion: n at f=0 equals 8

TC-FAT-008

Description: R=0.90, C=0.95, $\beta=2$, AF=1.0 $\rightarrow$ no β sensitivity table

Tests: UR-007

Command: `jrrun jrc_ss_fatigue.R 0.90 0.95 2 1.0`

Rationale: At AF = 1, $AF^{\beta} = 1$ so n is independent of β (code review 2026-10, COR-19)

Pass criterion: Exit 0; no 'Sensitivity to Weibull shape parameter' table; 'does not depend on the Weibull shape' note

10.6 Sample Size — Paired (jrc_ss_paired)

Formula: $n = \text{ceiling}(((z_{\alpha} + z_{\beta})/\text{effect_size})^2) + 1$. effect_size = delta/sd. 2-sided: $z_{\alpha} = \text{qnorm}(0.975)$. 1-sided: $z_{\alpha} = \text{qnorm}(0.95)$.

TC-PAIRED-006

Description: delta=0.5, sd=1.0, sides=2, power=0.90, C=0.95 $\rightarrow$ n=44

Tests: UR-008

Command: `jrrun jrc_ss_paired.R 0.5 1.0 2`

Rationale: $z_{\alpha} = \text{qnorm}(0.975) = 1.9600$; $z_{\beta} = \text{qnorm}(0.90) = 1.2816$; $n = \text{ceiling}((3.2416/0.5)^2) + 1 = \text{ceiling}(42.03) + 1 = 43 + 1 = 44$

Pass criterion: n at power=0.90, C=0.95 equals 44

TC-PAIRED-007

Description: delta=0.5, sd=1.0, sides=1, power=0.90, C=0.95 $\rightarrow$ n=36

Tests: UR-008

Command: `jrrun jrc_ss_paired.R 0.5 1.0 1`

Rationale: $z_{\alpha} = \text{qnorm}(0.95) = 1.6449$; $z_{\beta} = \text{qnorm}(0.90) = 1.2816$; $n = \text{ceiling}((2.9265/0.5)^2) + 1 = \text{ceiling}(34.26) + 1 = 35 + 1 = 36$

Pass criterion: n at power=0.90, C=0.95 equals 36 (< 44 from TC-PAIRED-006, as expected)

10.7 Sample Size — Equivalence (jrc_ss_equivalence)

TOST always uses 1-sided $z_{\alpha} = qnorm(C)$. For 2-sided equivalence both one-sided tests must reject, so β is split over the two sides: $z_{\beta} = qnorm(1 - (1 - power)/2)$. (Corrected in the 2026-10 code review, COR-10; previously $z_{\beta} = qnorm(power)$.)

TC-EQUIV-005

Description: delta=0.5, sd=1.0, sides=2, power=0.90, C=0.95 → n=45

Tests: UR-009

Command: jrrun jrc_ss_equivalence.R 0.5 1.0 2

Rationale: $z_{\alpha} = qnorm(0.95) = 1.6449$; $z_{\beta} = qnorm(0.95) = 1.6449$; $n = \text{ceiling}((3.2898/0.5)^2) + 1 = \text{ceiling}(43.29) + 1 = 44 + 1 = 45$

Pass criterion: n at power=0.90, C=0.95 equals 45

TC-EQUIV-006

Description: delta=0.5, sd=1.0, sides=2, power=0.95, C=0.95 → n=53

Tests: UR-009

Command: jrrun jrc_ss_equivalence.R 0.5 1.0 2

Rationale: $z_{\alpha} = qnorm(0.95) = 1.6449$; $z_{\beta} = qnorm(0.975) = 1.9600$; $n = \text{ceiling}((3.6049/0.5)^2) + 1 = \text{ceiling}(51.98) + 1 = 52 + 1 = 53$

Pass criterion: n at power=0.95, C=0.95 equals 53 (> 45 from TC-EQUIV-005, as expected)

10.8 Tolerance Interval Verification (jrc_verify_attr)

Dataset verify_attr_known.csv: mean=10.0000, sd=1.0000. $K1(30, 0.95, 0.95) = 2.220$; $K2(30, 0.95, 0.95) = 2.555$.

1-sided TL = $10.000 - 2.220 = 7.780$. 2-sided LTL = 7.445, UTL = 12.555.

TC-VER-009

Description: verify_attr_known.csv, P=0.95, C=0.95, LSL=7.0 → TL≈7.780, verdict

Tests: UR-017

Command: jrrun jrc_verify_attr.R 0.95 0.95 oq/data/verify_attr_known.csv value 7.0 -

Rationale: TL = mean - $K1 \times sd = 10.000 - 2.220 \times 1.000 = 7.780$; $7.780 > 7.0$ →

Pass criterion: Extracted 1-sided lower TL ≈ 7.780 ± 0.050 AND in output

TC-VER-010

Description: verify_attr_known.csv, P=0.95, C=0.95, LSL=8.0 → TL≈7.780 < 8.0, verdict

Tests: UR-017

Command: jrrun jrc_verify_attr.R 0.95 0.95 oq/data/verify_attr_known.csv value 8.0 -

Rationale: Same TL = 7.780; $7.780 < 8.0 = LSL$ →

Pass criterion: Extracted 1-sided lower TL ≈ 7.780 ± 0.050 AND TL < 8.0 AND in output

TC-VER-011

Description: verify_attr_known.csv, P=0.95, C=0.95, LSL=7.0, USL=13.0 → LTL≈7.445, UTL≈12.555, verdict

Tests: UR-017

Command: jrrun jrc_verify_attr.R 0.95 0.95 oq/data/verify_attr_known.csv value 7.0 13.0

Rationale: LTL= $10.000 - 2.555 = 7.445 > 7.0$; UTL= $10.000 + 2.555 = 12.555 < 13.0$ →

Pass criterion: 2-sided LTL ≈ 7.445 ± 0.050 AND UTL ≈ 12.555 ± 0.050 AND in output

10.9 Discrete Verification (jrc_verify_discrete)

Clopper-Pearson upper one-sided CI: upper_bound = $qbeta(C, f+1, N-f)$. For $f=0$: exact formula upper_bound = $1 - (1-C)^{1/N}$. PASS if upper_bound < $1-P$; FAIL otherwise. Margin = $(1-P) - \text{upper_bound}$.

TC-VER-DISC-001

Description: N=125, f=2, P=0.95, C=0.95 → exit 0,  in output

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 125 2 0.95 0.95

Pass criterion: exit code 0 AND  in output

TC-VER-DISC-002

Description: N=30, f=3, P=0.95, C=0.95 → exit 0,  in output

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 30 3 0.95 0.95

Pass criterion: exit code 0 AND  in output

TC-VER-DISC-003

Description: f=0 → exit 0, note directing user to jrc_ss_discrete_ci

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 60 0 0.95 0.95

Pass criterion: exit code 0 AND 'jrc_ss_discrete_ci' in output

TC-VER-DISC-004

Description: N=125, f=2: upper CI bound matches pure-Python Clopper-Pearson $\approx 0.04951 \pm 0.0005$

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 125 2 0.95 0.95

Rationale: Bisection: find u s.t. $\sum_{k=0}^2 C(125, k) \cdot u^k \cdot (1-u)^{(125-k)} = 0.05$; $u \approx 0.04951$

Pass criterion: Extracted upper bound (% / 100) within ± 0.0005 of $_cp_upper(125, 2, 0.95)$ computed by pure-Python bisection on Binomial CDF — independently of R

TC-VER-DISC-005

Description: N=60, f=0: upper CI bound matches exact formula $1-(0.05)^{(1/60)} \approx 0.04870 \pm 0.0005$

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 60 0 0.95 0.95

Rationale: Exact: $1-\exp(\ln(0.05)/60) = 1-\exp(-0.04993) \approx 0.04870$; no R computation involved

Pass criterion: Extracted upper bound (% / 100) within ± 0.0005 of $1-(0.05)^{(1/60)} \approx 0.04870$

TC-VER-DISC-006

Description: $f > N \rightarrow$ non-zero exit

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 10 15 0.95 0.95

Pass criterion: exit code $\neq 0$

TC-VER-DISC-007

Description: Only 3 arguments → non-zero exit, 'Usage' in output

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 125 2 0.95

Pass criterion: exit code $\neq 0$ AND 'usage' in output (case-insensitive)

TC-VER-DISC-008

Description: Margin positive for PASS case, negative for FAIL case

Tests: UR-025

Command: jrrun jrc_verify_discrete.R 125 2 0.95 0.95 (PASS)
jrrun jrc_verify_discrete.R 30 3 0.95 0.95 (FAIL)

Pass criterion: Extracted margin > 0 for PASS case; extracted margin < 0 for FAIL case

11. Requirements Traceability Matrix

All 164 test cases in the community OQ suite (oq/), mapped to user requirements. Phase 9 numeric TCs are highlighted.

UR	Test Cases	Count	Test File
UR-001	TC-DISC-001..009	9	test_ss_suite.py
UR-002	TC-DISCICI-001..006	6	test_ss_suite.py
UR-003	TC-ATTR-001..010	10	test_ss_suite.py
UR-004	TC-ATTRCK-001..006	6	test_ss_suite.py
UR-005	TC-ATTRCI-001..005	5	test_ss_suite.py
UR-006	TC-SIGMA-001..007	7	test_ss_suite.py
UR-007	TC-FAT-001..008	8	test_ss_suite.py
UR-008	TC-PAIRED-001..007	7	test_ss_suite.py
UR-009	TC-EQUIV-001..006	6	test_ss_suite.py
UR-010	TC-GRR-001..005	5	test_ss_suite.py
UR-011	TC-NORM-001..006	6	test_diagnostic_suite.py
UR-012	TC-OUT-001..005	5	test_diagnostic_suite.py
UR-013	TC-CAP-001..004	4	test_diagnostic_suite.py
UR-014	TC-DESC-001..004	4	test_diagnostic_suite.py
UR-015	TC-BA-001..006	6	test_statistical_suite.py
UR-016	TC-WEIB-001..006	6	test_statistical_suite.py
UR-017	TC-VER-001..011	11	test_statistical_suite.py
UR-018	TC-GEN-* (5 scripts × ~4)	21	test_gen_suite.py
UR-019	TC-CCSV-001..008	8	test_convert_suite.py
UR-020	TC-CTXT-001..007	7	test_convert_suite.py
UR-021	Error-path TCs (all suites)	—	all test files
UR-022	TC-CORE-OQ-004	1	test_core.py
UR-023	TC-DOE-DES-001..019, TC-DOE-ANA-001..018	37	test_doe_suite.py
UR-024	TC-CORE-IQ-001, TC-CORE-OQ-001..005	6	test_core.py
UR-025	TC-VER-DISC-001..008	8	test_statistical_suite.py
Total	All community OQ tests	198	oq/

12. Deferred Test Cases

The following test cases are planned but deferred to Phase 10 (future validation cycle):

- Non-normal data path for jrc_ss_attr and jrc_verify_attr (Box-Cox transformation): numeric assertions for skewed datasets where the script applies Box-Cox before computing the K-factor. The existing functional tests (TC-ATTR-004, TC-VER-004) verify the code path is exercised but do not assert specific transformed values.
- Gauge R&R numeric assertions (TC-GRR-005+): numeric %GRR and ndc values for jrc_msa_grr_design.